

# Demystifying HPC

How to Use and Optimize for HPC Systems

Chris Stylianou

Research Engineer

CaSToRC

[eurocc.cyi.ac.cy](http://eurocc.cyi.ac.cy)

- Research Engineer at CaSToRC, PhD
- *Training and Skills Development* Task Leader for EuroCC2
- Area of expertise in High Performance Computing (HPC)
- Contact & Info:
  - Email: [c.stylianou@cyi.ac.cy](mailto:c.stylianou@cyi.ac.cy)
  - Website: [cstyl.github.io](https://cstyl.github.io)




You can find the slides at: <https://github.com/CaSToRC-Cyl/EuroCC-HPC-with-Python-2025>



## Accelerated Simulations

Solve complex simulations faster by leveraging parallel processing



## Enhanced Model Complexity

Simulations with higher spatial and temporal resolution  
Train large-scale ML models (Bln/Tln Parameters)



## Ability to Tackle Larger Problems

Scale to larger domain sizes that are infeasible on local machines  
Process massive datasets



## Cost-Effectiveness

Reduce the need for physical experiments by conducting virtual testing  
Reduced need to invest in high-end computing infrastructure.

## Supercomputers

### Specialized Hardware:

- Optimized for **parallel** and **capability computing**
- solving the **largest** and most **complex problems**
- High-speed interconnects.

### Purpose:

- Focused on scientific research, simulations, and high-performance tasks.

### Resource Allocation:

- Fixed and highly controlled environment.
- Users **share resources** via job **schedulers** like SLURM.

## Cloud Computing

### General-Purpose Hardware:

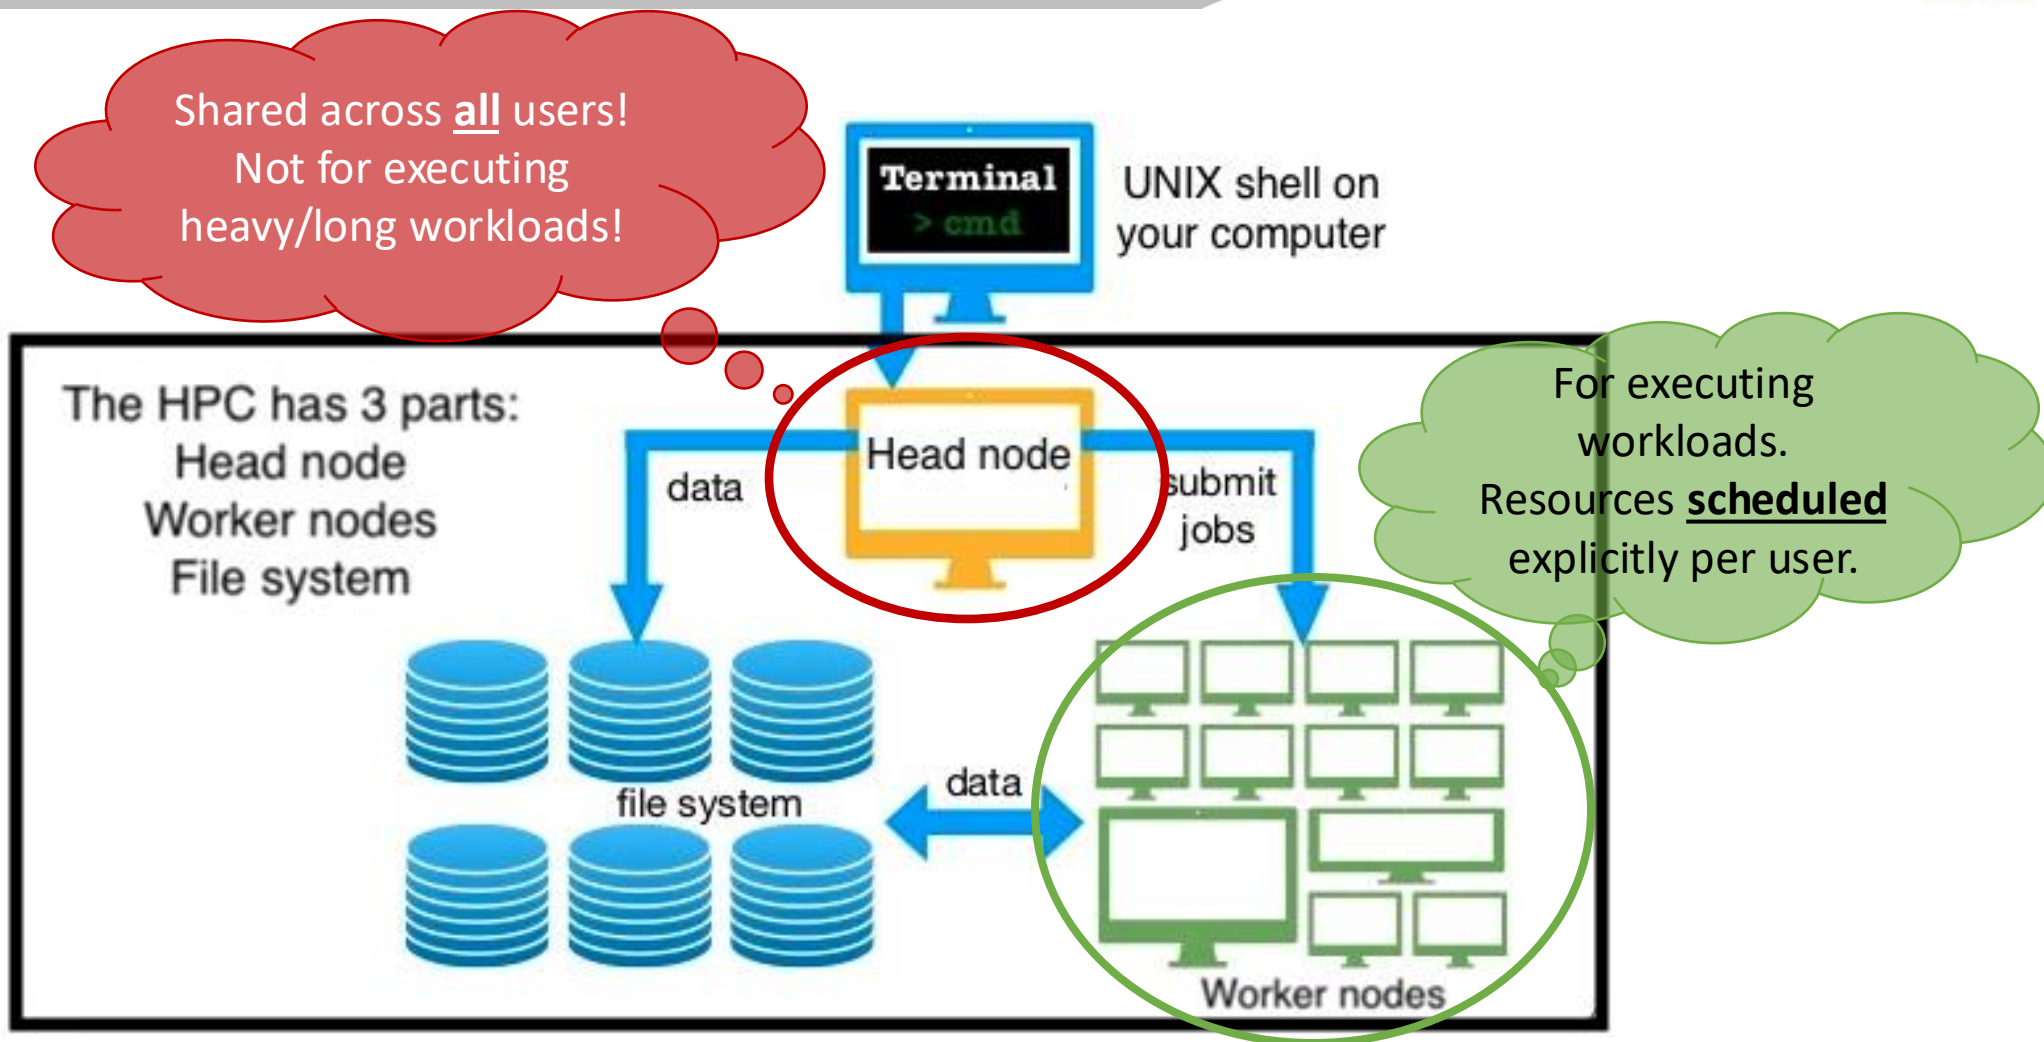
- Flexible and scalable **virtual machines**.
- Designed for **capacity computing**
  - handling many smaller tasks concurrently.
- Standard networking and storage infrastructure.

### Purpose:

- Ideal for diverse workloads, including web applications, and general-purpose computing.

### Resource Allocation:

- On-demand provisioning with **pay-as-you-go pricing**.
- Easily scales up or down based on need.





## Access and Setup

Secure access with credentials (e.g., SSH keys).  
Understand system layout (home, scratch, project directories).



## Command-Line Basics

Navigate and manage files using Linux commands.  
Edit scripts with tools like vim, nano or Emacs.



## Job Submission

Write job scripts specifying CPUs, GPUs, and memory.  
Submit (sbatch) and monitor (squeue) jobs with a scheduler.



## Software Modules

Load required software (module load).



## Data Management

Transfer files efficiently (e.g., scp, rsync).  
Organize inputs/outputs and monitor storage usage.

## How to Access Cyclone

- Use **Secure Shell (SSH)** to log in to the supercomputing system.
- Requires a **username** and **password** or **key-based authentication**.

## Tools for SSH by Platform:

- **Linux/Mac:** SSH is pre-installed; use the terminal.
- **Windows:**
  - Use tools like **PuTTY** or **Windows Subsystem for Linux (WSL)** for SSH access.
  - **Tip:** Install OpenSSH through PowerShell or the Settings menu.
    - Requires **Administrative privileges!**

## • Example SSH Command:

```
$ ssh -i /path/to/ssh/private/key cstyl@cyclone.hpcf.cyi.ac.cy
```





NONE of the data on this system is backed up. Please keep backups of valuable files at alternative locations. Maintaining backups is the sole responsibility of the user.

Welcome to the Cyclone HPC system!

- System Documentation: <https://hpcf.cyi.ac.cy/documentation/index.html>
- Support: Email to [hpc.support@cyi.ac.cy](mailto:hpc.support@cyi.ac.cy) or open a ticket here: <https://hpcf-support.atlassian.net/servicedesk/customer/portal/3>
- Use "module avail" and "module load" to list/use available software
- You can now use the "qhist" command to check your project's usage
- Utilize --reservation=short for 1-hour jobs on specific nodes during work hours for no queue time

```
Last login: Thu Nov 21 16:06:56 2024 from 46.199.227.111
(base) [cstyl@front02 ~]$
```



## Key Linux Commands:

Command	Purpose	Example
<b>ls</b>	List directory contents	<code>ls -l</code>
<b>cd</b>	Change directory	<code>cd /path/to/directory</code>
<b>pwd</b>	Print the current directory path	<code>pwd</code>
<b>cp</b>	Copy files or directories	<code>cp file1 file2</code>
<b>mv</b>	Move or rename files or directories	<code>mv oldname newname</code>
<b>rm</b>	Remove files (use cautiously!)	<code>rm file</code>

## Best Practices:

- Use `ls` and `pwd` to verify your current location.
- Avoid running `rm` without confirming the file path.



```
(base) [cstyl@front02 ~]$ pwd
```

```
/nvme/h/cstyl
```

```
(base) [cstyl@front02 ~]$ ls -l
```

```
total 3783350
```

lrwxrwxrwx 1 cstyl qcd	15 Jun 14 2022	<a href="#">data_p069</a> -> <a href="#">/onyx/data/p069</a>
lrwxrwxrwx 1 cstyl qcd	15 Jul 10 2023	<a href="#">data_p163</a> -> <a href="#">/onyx/data/p163</a>
lrwxrwxrwx 1 cstyl qcd	15 Jul 11 2023	<a href="#">data_p165</a> -> <a href="#">/onyx/data/p165</a>
lrwxrwxrwx 1 cstyl qcd	20 Apr 17 2024	<a href="#">edu20</a> -> <a href="#">/nvme/scratch/edu20/</a>
drwxr-xr-x 4 cstyl qcd	9 Apr 27 2024	<a href="#">gpt-fast</a>
lrwxrwxrwx 1 root root	20 Mar 11 2024	<a href="#">scratch</a> -> <a href="#">/nvme/scratch/cstyl/</a>
drwxr-xr-x 13 cstyl qcd	15 Apr 16 2024	<a href="#">wee_archie</a>



## Directory Overview

**Home Directory:** Persistent storage for scripts and small files.

`$HOME=/nvme/h/<username>`

**Scratch Directory:** Temporary, high-speed storage for job data.

`$SCRATCH=/nvme/scratch/<username>`

**Shared Directory:** Collaboration space for team projects.

`$DATA_<pid>=/onyx/data/<pid>`



## Best Practices for File Management

Store **active jobs** in the **scratch directory**.

Store **source code** and **build executables** in **home directory**.

Store **large** shared project **data** in **shared directory**.

Move important results to **home** or external storage to prevent loss.

**Note: NO BACKUPS**

## What Are Modules?

- **Modules** manage software environments, ensuring compatibility with HPC resources.
- Load, switch, or unload software dynamically.
- **Note: No sudo access on HPC Systems!**

### Key Commands:

Command	Purpose	Example
<b>module avail</b>	List all available software modules	module avail
<b>module load</b>	Load a specific software module	module load gcc
<b>module unload</b>	Unload a specific module	module unload gcc
<b>module list</b>	List currently loaded modules	module list

**Tip:** Check module dependencies to avoid conflicts.

```
(base) [cstyl@front02 ~]$ python --version
```

```
Python 3.10.13
```

```
(base) [cstyl@front02 ~]$ module avail Python/
```

```
----- /eb/modules/all -----
```

```
Python/2.7.16-GCCcore-8.3.0
```

```
...
```

```
Python/3.10.8-GCCcore-12.2.0
```

```
Python/3.11.5-GCCcore-13.2.0
```

```
(base) [cstyl@front02 ~]$ module load Python/3.11.5-GCCcore-13.2.0
```

```
(base) [cstyl@front02 ~]$ python --version
```

```
Python 3.11.5
```



## Definition

**SLURM** (Simple Linux Utility for Resource Management) is a job scheduler that allocates resources and manages workloads on supercomputers.

It ensures efficient sharing of resources among multiple users.




## Key Roles

**Resource Allocation:** Assigns CPUs, memory, GPUs, and other resources to jobs.

**Job Scheduling:** Prioritizes jobs based on factors like resource availability and queue position.

**Monitoring:** Tracks the status of jobs and manages failures or cancellations.



## Why SLURM?

Maximizes system utilization.

Ensures **fair resource distribution**.

A SLURM script is a text file with commands that specify:

1. **Job Information:** Job name, output files, email notifications.
2. **Resource Requests:** CPUs, memory, GPUs, and wall time.
3. **Execution Commands:** The actual program or script to run.

```
1  #!/bin/bash
2  #SBATCH --job-name=example_job    # Job name
3  #SBATCH --output=output.txt       # Output file
4  #SBATCH --error=error.txt         # Error file
5  #SBATCH --ntasks=1               # Number of tasks (CPUs)
6  #SBATCH --mem=4G                 # Memory allocation
7  #SBATCH --time=01:00:00          # Time limit (hh:mm:ss)
8
9  module load Python               # Load required module
10
11 python script.py                 # Run Python script
```



## Example Workflow:

1. Submit a job using sbatch.
2. Monitor progress with squeue.
3. Cancel if necessary with scancel.

## Job Status Symbols:

- **R**: Running.
- **PD**: Pending (waiting in queue).
- **F**: Failed.

Category	Command	Action
Job Submission	<code>sbatch &lt;job_script&gt;</code>	Submit a batch job
	<code>srun &lt;command&gt;</code>	Run a command in parallel
Job Monitoring	<code>squeue</code>	View queued jobs
	<code>squeue --user=\$USER</code>	View MY queued jobs
	<code>scontrol show job &lt;job_id&gt;</code>	Detailed job info
Job Management	<code>scancel &lt;job_id&gt;</code>	Cancel a job
	<code>scontrol hold &lt;job_id&gt;</code>	Hold a job
	<code>scontrol release &lt;job_id&gt;</code>	Release a held job
Resource Allocation	<code>sinfo</code>	Information about nodes and partitions
	<code>scontrol show node &lt;node_id&gt;</code>	Show Node details

## Best Practices:

- Request only what your job needs to avoid wasting resources.
- Use `--time` to limit job runtime and prioritize efficiency.
- Use `--account` to specify which project to “charge”.

Resource	Description	Job Specification
Nodes	Number of Nodes	<code>--nodes</code>
CPU tasks	Number of CPU tasks per node	<code>--ntasks-per-node</code>
CPU threads	Number of CPU threads (cores) per task	<code>--cpus-per-task</code>
Memory	Amount of RAM per Job	<code>--mem</code>
GPUs	Request for GPUs if needed	<code>--gres</code>
System Partition	Either use the CPU or GPU part of the system	<code>--partition</code>

## Command Workflow:

1. Write the submission script (python\_job.sh)
2. Submit: **sbatch** python\_job.sh
3. Monitor: **squeue** -u <username>
4. Cancel (if needed): **scancel** job\_id

## Output:

- Job results and logs will appear in python\_output.txt, errors in python\_error.txt.
- Check logs for troubleshooting if the job fails.

```
1  #!/bin/bash
2  #SBATCH --job-name=python_job
3  #SBATCH --output=python_output.txt
4  #SBATCH --error=python_error.txt
5
6  #SBATCH --partition=cpu
7  #SBATCH --nodes=1
8  #SBATCH --ntasks-per-node=1
9  #SBATCH --cpus-per-task=1
10 #SBATCH --mem=2G
11 #SBATCH --time=00:30:00
12
13 module load Python
14 python example_script.py
```

```
[(base) [cstyl@front02 ~]$ squeue
```

JOBID	PARTITION	NAME	USER	ST	TIME	NODES	NODELIST(Reason)
998469_[19-25%1]	cpu	CNO_300K	sbhowmic	PD	0:00	1	(JobArrayTaskLimit)
1007611_[4-5%1]	cpu	Ag1H+_D1	sbhowmic	PD	0:00	1	(JobArrayTaskLimit)
1008164	cpu	enthalpy	cy22kp1	PD	0:00	2	(Dependency)
1008163	cpu	analysis	cy22kp1	PD	0:00	1	(Dependency)
1008403_[141-146]	cpu	nanli44c	cy24nl1	PD	0:00	1	(Priority)
1008402_[131-136]	cpu	nanli44c	cy24nl1	PD	0:00	1	(Priority)
1008392_[66-68]	cpu	nanli44c	cy24nl1	PD	0:00	1	(Resources)
998469_18	cpu	CNO_300K	sbhowmic	R	15:45:59	1	cn09
1007611_3	cpu	Ag1H+_D1	sbhowmic	R	10:08:58	1	cn02
1008334	cpu	data_min	jmoreno	R	23:43:31	1	cn06
1008525	cpu	zn	eg20ra1	R	17:30:39	2	cn[15-16]
1008392_65	cpu	nanli44c	cy24nl1	R	7:51:38	1	cn13
1008392_63	cpu	nanli44c	cy24nl1	R	14:24:14	1	cn05
1008392_64	cpu	nanli44c	cy24nl1	R	14:24:14	1	cn12
1008392_61	cpu	nanli44c	cy24nl1	R	17:12:05	1	cn14
1008392_62	cpu	nanli44c	cy24nl1	R	17:12:05	1	cn17
1008392_59	cpu	nanli44c	cy24nl1	R	17:32:16	1	cn03
1008392_60	cpu	nanli44c	cy24nl1	R	17:32:16	1	cn04

## Why Use Alternative Tools?

- Enhance **productivity** with **user-friendly interfaces**.
- **Simplify file management** and script editing for supercomputing workflows.
- Provide advanced features like **SSH integration**, **file transfer**, and **remote execution**.

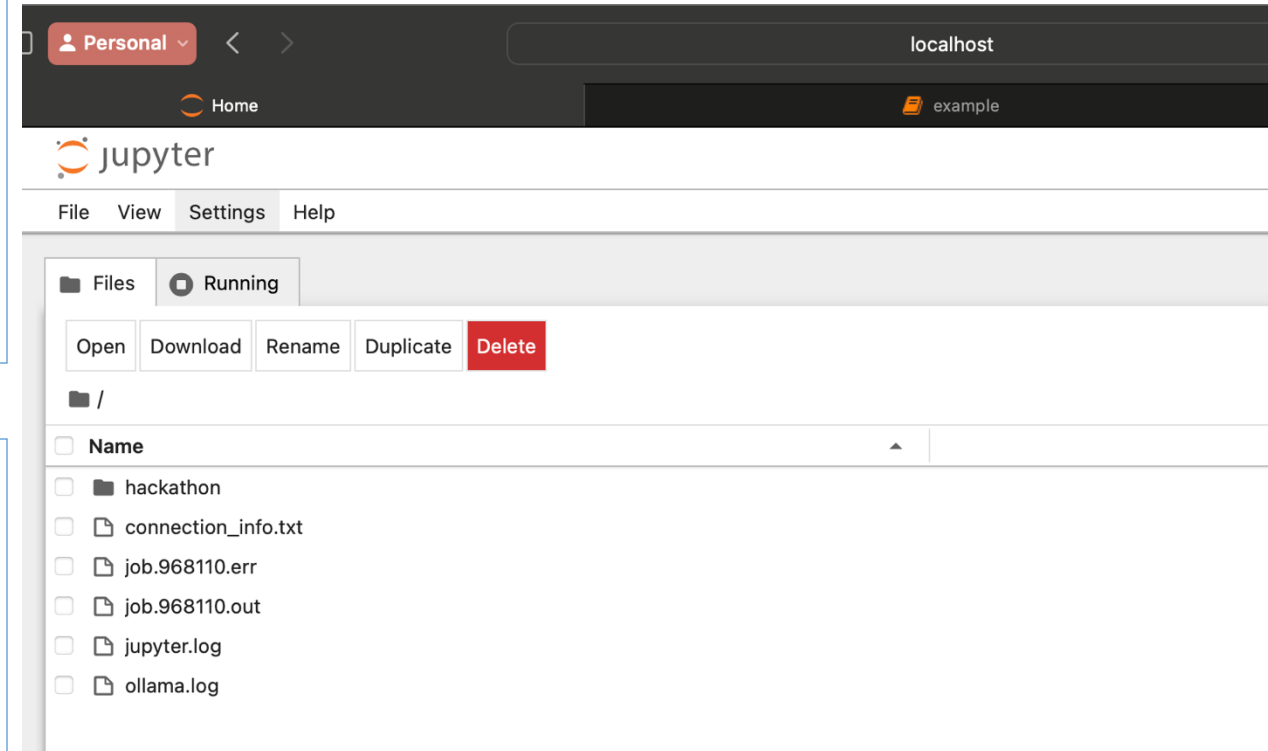
Tool	Purpose	Best For	Platform
<b>VS Code</b>	Remote editing and debugging	Writing scripts, interactive debugging, SSH access	<b>Windows, Mac, Linux</b>
<b>MobaXTerm</b>	SSH client with GUI	SSH access, file management	Windows
<b>FileZilla</b>	File transfer	Transferring data to/from supercomputers	Windows, Mac, Linux

## Advantages of Jupyter Notebooks:

- **Interactive Workflows:** Combines code, visualizations, and narratives in a single interface.
- **AI and Data Analysis:** Perfect for tasks like data exploration, visualization, and model development.
- **Ease of Use:** Intuitive, web-based platform that lowers the learning curve for complex workflows.

## Why Use HPC for Jupyter?

- Access to powerful compute resources (e.g., CPUs, GPUs).
- Handle large-scale datasets beyond the capabilities of local machines.
- Perform **AI training** and **numerical simulations** interactively and efficiently.





## Maximum Wall Time for Jobs

Jobs are subject to a maximum runtime (e.g., 24 hours on Cyclone).

Long-running tasks may need checkpointing or splitting.



## Compiler/Software Versions

Availability is restricted to pre-installed compilers and software.

Users must adapt to available versions or request installations.



## No Sudo/Root Access

Administrative privileges are not granted.

Custom software or dependencies must be compiled within user space.



## Queue Wait Times

Jobs may experience delays due to high demand for resources.

Resource availability depends on job priority and system load.



## Shared Resources

Performance may vary due to shared compute and storage infrastructure.

Careful resource allocation is required to avoid bottlenecks.



## Start with Serial Execution

- Begin with a single-threaded version of your application - Verify correctness and establish a performance baseline.
- **Development Time: Low** – Minimal effort to get a working version.




## Transition to Multi-Threading (Single Node)

- Use parallel programming models like OpenMP for **shared memory** - Optimize for **multi-core CPUs** on a single node.
- **Development Time: Moderate** – Requires identifying parallel regions and potential thread synchronization.



## Scale to Multiple Nodes (Distributed Computing)

- Implement distributed memory parallelism using MPI. - Partition workloads and minimize inter-node communication.
- **Development Time: High** – Significant effort to design communication patterns and debug distributed code.



### Leverage Heterogeneous Architectures (GPUs)

- **Offload compute-intensive tasks** to GPUs using CUDA, OpenACC, or HIP. - Scale further with multi-GPU setups across nodes.
- Optimize **data transfers** between host (CPU) and device (GPU).
- **Development Time: Very High** – Requires specialized knowledge of GPU programming and iterative tuning.



### Profile and Optimize at Each Stage

- Profile applications after every major change to identify resource bottlenecks (CPU, memory, I/O, communication).
- Use tools like gprof, perf, or GPU-specific profilers (e.g., nvprof, Nsight).
- **Development Time: Ongoing** – Profiling and optimization are iterative and critical for scalability.



## Problem Definition

- Define Governing Equations and set up the computational domain and boundary conditions

## Discretization

- Break down the continuous domain into discrete grid.
- FDM, FEM, FVM

## Solver Setup

- Define Numerical Schemes for solving equations and select Solvers

## Preprocessing

- Generate Computational Mesh, Initialise Parameters and Boundary Conditions

## Solution Process

- Execute Solver – Iteratively solve equations and monitor convergence

## Post Processing

- Visualisation of Results



## Understand the Serial Application

Profile the code to identify bottlenecks.

Verify scalability potential for parallelization.



## Prepare for Parallelization

Refactor code for modularity.

Choose parallel paradigms: OpenMP (shared), MPI (distributed), CUDA/GPU.



## Parallelize the Code

**Single Node (OpenMP):**

Parallelize loops and hotspots.

**Multi-Node (MPI):**

Partition domain, optimize communication.

**GPU Acceleration:**

Offload compute-intensive tasks to GPUs.



## Test and Debug

Validate accuracy against the serial version.

Debug using parallel tools.



## Optimize for HPC Systems

Improve load balancing and memory usage.

Reduce communication overhead and use optimized libraries (e.g., PETSc)




## Scale the Application

Perform strong and weak scaling tests.

Adapt for multi-node and multi-GPU systems.





- Supercomputers and cloud systems serve different purposes (**capability computing** vs. **capacity computing**).
- Supercomputers allow us to solve larger and more complex problems beyond the capabilities of personal computers or workstations
- Efficient navigation, resource management, and job scheduling are essential for success.
- Tools like Jupyter and alternative IDEs enhance **usability** and **productivity**.
- **Scalability is a Journey**: Porting and scaling applications require a step-by-step approach
  - From serial execution to multi-node and heterogeneous computing.
- **Development Effort vs. Performance**: Scaling offers significant benefits but requires careful planning and expertise.



🏠 / Tutorials / Introduction to HPC Systems

<https://castorc-cyi.github.io/eurocc-tutorials/>

## 1. Introduction to HPC Systems

### 1.1. Overview

This tutorial provides a high-level introduction to High-Performance Computing (HPC) systems, with a focus on Cyclone's architecture and operational principles. Participants will explore the fundamental components of an HPC system, including compute nodes and file systems, while gaining insight into data management policies. Additionally, the tutorial introduces key concepts such as software management through modules, and job scheduling using SLURM, setting the stage for deeper exploration in subsequent tutorials.

### 1.2. Learning Objectives

By the end of this tutorial, participants will be able to:

1. Describe the architecture of an HPC system, including Cyclone's compute nodes and interconnects.
2. Identify and understand the use cases of Cyclone's file systems (home, scratch, shared directories).
3. Identify when to use an HPC system or alternative solutions such as Cloud systems or High-end Workstations.
4. Recognize the role of modules in managing software environments and how they simplify system use.
5. Understand the purpose of job scheduling and the function of SLURM in resource management.

### 1.3. Overview of HPC Architectures

An HPC system is typically composed of multiple interconnected compute nodes which work in





Thank you for your attention!



More information:



<https://eurocc.cyi.ac.cy/>  
<https://www.linkedin.com/company/eurocc2>



Contact us at:  
[eurocc-contact@cyi.ac.cy](mailto:eurocc-contact@cyi.ac.cy)



RESEARCH  
& INNOVATION  
FOUNDATION



**EuroHPC**  
Joint Undertaking

Funded by the European Union. This work has received funding from the European High Performance Computing Joint Undertaking (JU) and Germany, Bulgaria, Austria, Croatia, Cyprus (co-funded by the EU within the framework of the Cohesion Policy Programme “THALIA 2021-2027”), Czech Republic, Denmark, Estonia, Finland, Greece, Hungary, Ireland, Italy, Lithuania, Latvia, Poland, Portugal, Romania, Slovenia, Spain, Sweden, France, Netherlands, Belgium, Luxembourg, Slovakia, Norway, Türkiye, Republic of North Macedonia, Iceland, Montenegro, Serbia under grant agreement No 101101903.